

Open Resilient Cluster Manager: A Distributed Approach to a Resilient Router Manager

Ralph H. Castain, Ph.D.
Greenplum/EMC

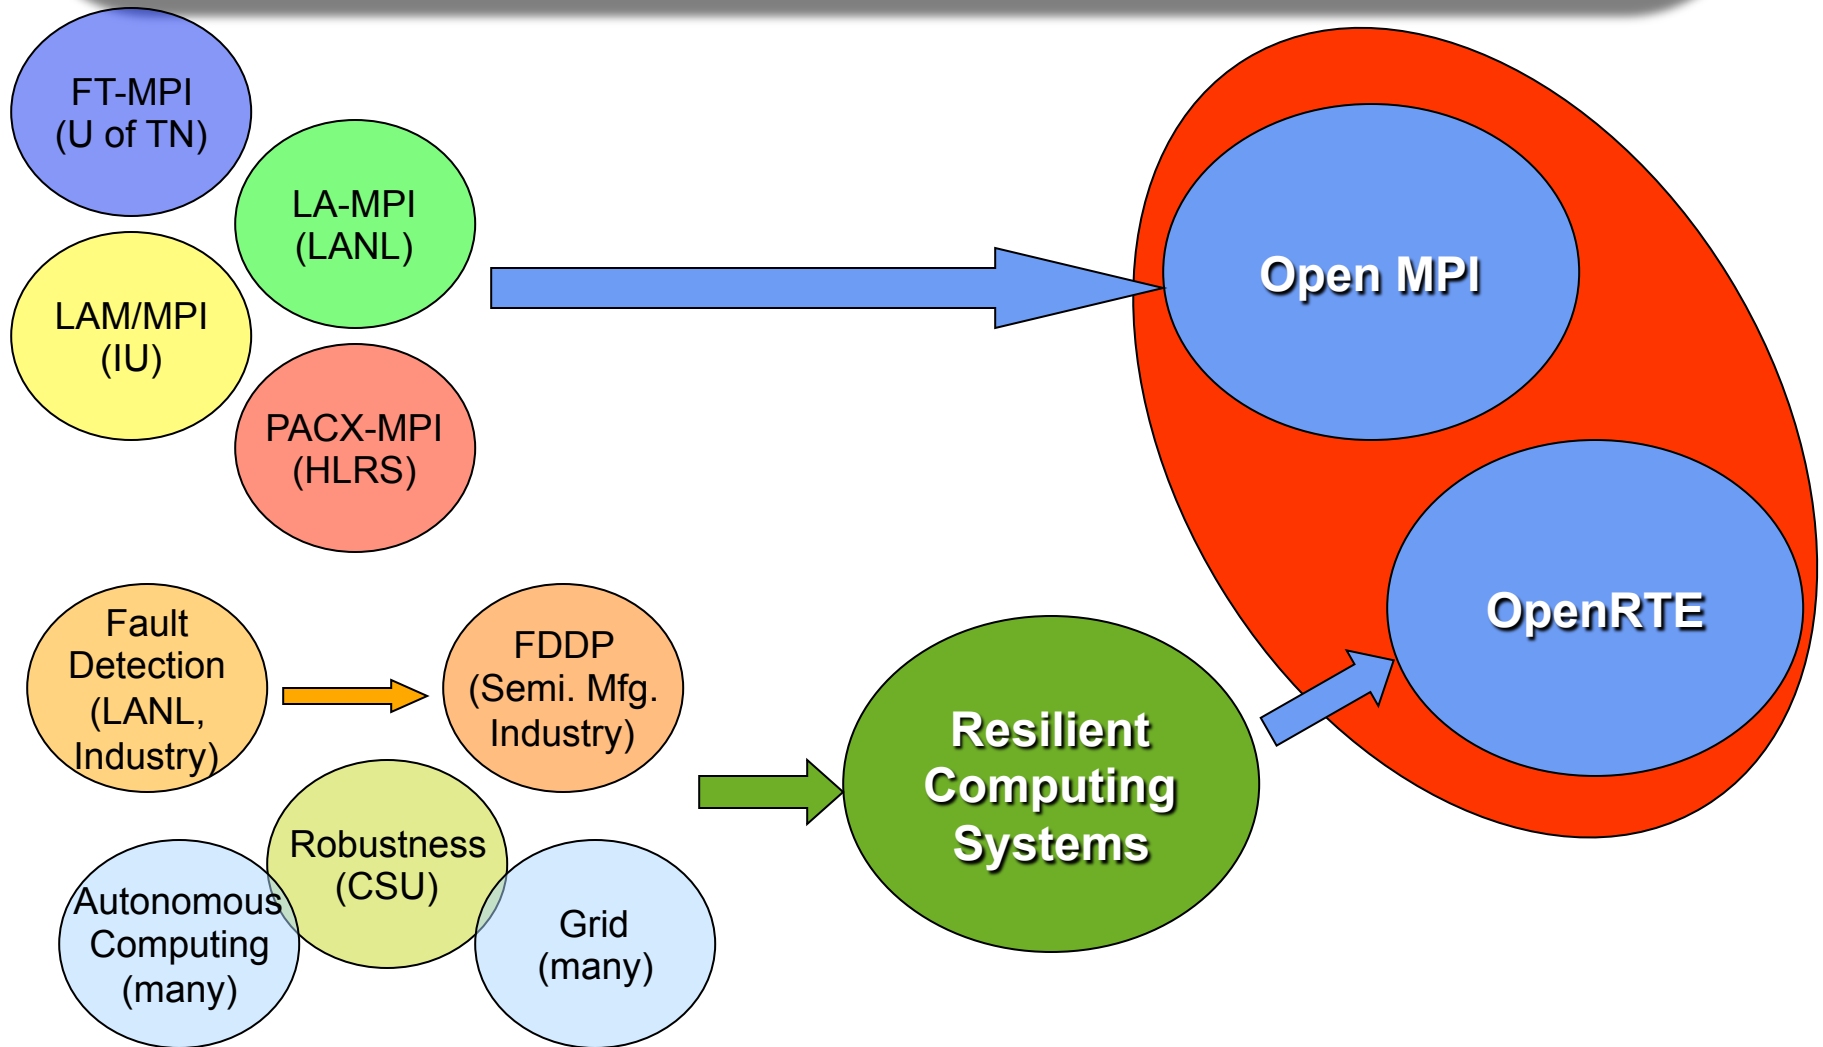
Objectives

- 100% uptime
 - Fully resilient to faults
 - On-the-fly maintenance
 - On-the-fly upgrades
- Test upgrades in the field without impact
 - Operate trial software systems in parallel, completely isolated
 - Transparent switch to production
- Maintain or enhance performance
 - Balance loads across all available processors

The Approach


- Distributed redundancy
 - NO master
 - Multiple copies of everything
 - Running in tracking mode
 - Parallel, seeing identical input
 - Multiple ways of selecting “leader”
- Utilize component architecture
 - Multiple ways to do something => plug-in!
 - Encourage experimentation
 - Build on OpenRTE

ORTE: Convergence of Ideas



Reliance on Plug-ins

- Formalized interfaces
 - Specifies “black box” implementation
 - Different implementations available at run-time
 - Can compose different systems on the fly



OpenRTE and Plug-ins

- Plug-ins are shared libraries
 - Central set of plug-ins in installation tree
 - Users can also have plug-ins under \$HOME
- Can add / remove plug-ins after install
 - No need to recompile / re-link apps
 - Download / install new plug-ins
 - Develop new plug-ins safely
- Update “on-the-fly”
 - Add, update plug-ins while running
 - Frameworks “pause” during update

Plug-in Benefits

- Stable, production quality environment for R&D
 - Can experiment inside the system without rebuilding everything else
 - Small learning curve (learn a few plug-ins, not the entire implementation)
 - Allow wide use, experience before exposing work
- Quickly roll out support for new platforms, features
 - Write only the plug-ins you want/need to change
 - Protect intellectual property (binary distro)

ORTE: Resiliency

- Fault
 - ***Events that hinder the correct operation of a process.***
 - May not actually be a “failure” of a component, but can cause system-level failure or performance degradation below specified level
 - Effect may be immediate or some time in the future.
 - Usually are rare. May not have many data examples.
- Fault prediction
 - Estimate probability of incipient fault within some time period in the future
- Fault Tolerance*reactive, static*
 - Ability to recover from a fault
- Robustness.....*metric*
 - How much can the system absorb without catastrophic consequences
- Resilience.....*proactive, dynamic*
 - Dynamically configure system to minimize impact of potential faults

Key Frameworks

Error Manager (Errmgr)

- Receives all process state updates
 - Sensor, waitpid
 - Includes predictions
- Determines response strategy
 - Restart locally, globally, abort
- Executes recovery
 - Accounts for fault groups to avoid repeated failover

Sensor

- Monitors software and hardware state-of-health
 - Sentinel file size, mod & access times
 - Memory footprint
 - Temperature
 - Heartbeat
 - ECC errors
- Predicts incipient faults
 - Trend, fingerprint
 - AI-based algos coming

Universal PNP

- Widely adopted standard
- ORCM uses only a part
 - PNP discovery via announcement on common multicast group
 - Includes application id, contact info
 - All applications respond
 - Wireup “storm” limits scalability
 - Various algorithms for storm reduction
 - APIs provided for sending/recving data
 - Broadcast-like output via multicast to anyone registered to recv output from that application
 - Unicast for direct IPC


Multicast: UDP

- Recvr-based methodology
 - All output sent to output group
 - Apps that want output from some other app join that app's output multicast group
- Each process joins several multicast groups
 - Given input, output for that app
 - System (announcements)
 - Error (recv process failure announcements)
- NACK-oriented protocol
 - Continuous stream of messages
 - Header contains sequence number
 - Missed message recovered via direct unicast from sender

Multicast: TCP

- Sender-based methodology
 - Output unicast directly to registered subscribers only
 - Each process maintains its own list of subscribers to each channel
 - Given contact info for daemon at startup
- Daemons
 - Simulate broadcast of all standard system channel messages (announcements, failures) by routing them across peers to apps
- Pros/cons
 - Resolves UDP messaging limits, lost message issues
 - Takes work to deal with things like race conditions on failure of subscribers

ORCM DVM



- One daemon (orcmond) per node
 - Started at node boot or launched by tool
 - Locally spawns and monitors processes, local system health sensors
 - Small footprint ($\leq 1\text{Mb}$)
- Each daemon tracks existence of others, has complete picture of system
 - PNP wireup
 - Know where all processes are located

orcm-sched

- Receives configuration input
 - Config mgr limitation => only one scheduler
 - Two mapping options
 - Maps process placement across available nodes, sends map to daemons
 - Sends configuration to daemons, daemons self-determine local procs
 - Daemons start local processes as required
- Provides point-of-contact for operational data requests
 - Returns snapshot of overall system state
- Responds to process failure (mapping mode 1)
 - Remaps process to new location

orcm-sched Failure

- Automatically restarted
 - Down time: ~2msec
 - No operational data available during down time
- Recover existing state from daemons upon startup
 - Each daemon provides current state
 - Scheduler assembles global picture
 - Total time to recover: ~100msec
- Receives configuration
 - Compares to existing state
 - Executes required corrections to match configuration

Multiple Parallel DVMs

- Allows
 - Concurrent development, testing in production environment
 - Sharing of development resources
- Unique identifier (ORTE jobid)
 - Maintains separation between orcmd's
 - Each application belongs to their respective DVM
 - No cross-DVM communication allowed


Apps: Multiple Replicas

- Multiple copies of each executable
 - Run on separate fault groups
 - Async, independent
- Shared pnp channels
 - Input: recvd by all
 - Output: broadcast to all, recvd by those who registered for input
 - Orcm-pnp delivers data to app layer only from leader => filtered so app only receives a single copy of each data stream

Leader Selection

- Two forms of leader selection
 - Internal to ORCM DVM
 - External facing
- Internal – framework, multiple plug-ins
 - App-specific plug-in
 - Configuration specified leader
 - Lowest rank alive
 - First contact
 - Leader-less (app receives raw data streams)

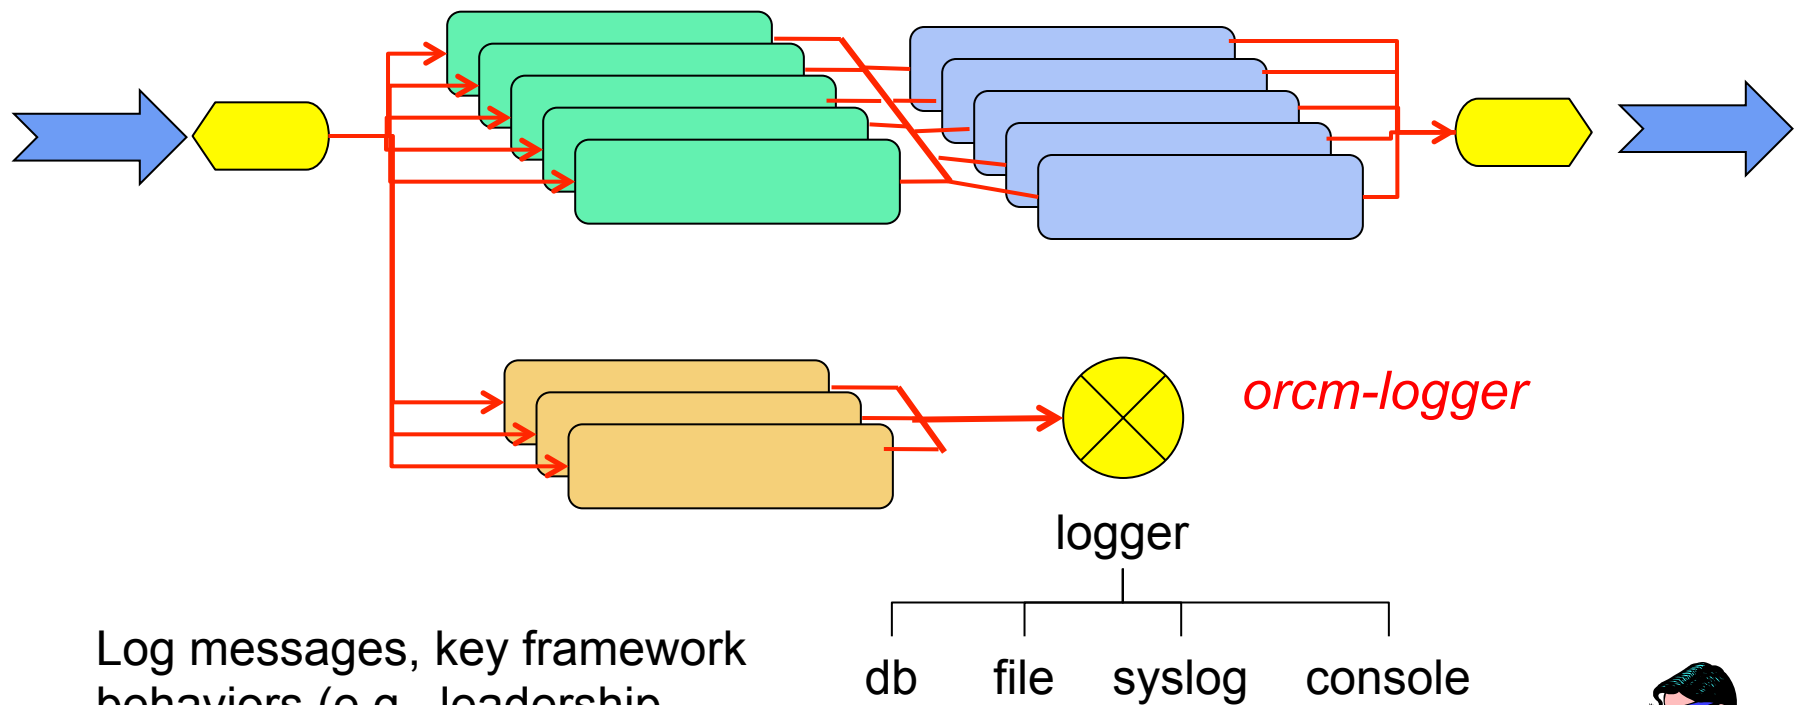
External Connections



orcm-connector

- Input
 - Broadcast on respective PNP channel
- Output
 - Determines “leader” to supply output to rest of world
 - Utilize any leader method in framework

Testing in Production



Log messages, key framework behaviors (e.g., leadership selections)



Software Maintenance

- On-the-fly plug-in activation
 - Configuration manager can select new plug-ins to load, reload, activate
 - Change priorities of active plug-ins
- Full replacement
 - When more than a plug-in needs updating
 - Start replacement version, test until validated
 - Configuration manager switches “leader”
 - Stop old version

Detecting Failures

- Application failures - detected by local daemon
 - Monitors for self-induced problems
 - Memory and cpu usage
 - Orders termination if limits exceeded or are trending to exceed
 - Detects unexpected failures via waitpid
- Hardware failures
 - Local hardware sensors continuously report status
 - Read by local daemon
 - Projects potential failure modes to pre-order relocation of processes, shutdown node
 - Detected by DVM when daemon misses heartbeats

Resilient Mapper

- Fault groups
 - Nodes with common failure mode
 - Node can belong to multiple fault groups
 - Defined in system file
- Map instances across fault groups
 - Minimize probability of cascading failures
 - One instance/fault group
 - Pick lightest loaded node in group
 - Randomly map extras
- Next-generation algorithms
 - Failure mode probability => fault group selection

Application Fault Response

- Local daemon
 - Detects (or predicts) failure
 - Notifies all running processes of failure
 - New leaders selected as required
- orcm-sched
 - Utilizes resilient mapper to determine re-location
 - Sends launch message to all daemons
- Replacement app
 - Announces itself on application public address channel
 - Receives responses – wires up to recv input as specified by app

Application State Recovery

- Apps register callback at startup
 - Indicate ability to provide state recovery data for peers
- On startup
 - Apps provided with “incarnation” number
 - If restart detected, apps call orcm API to request state recovery data
 - Request sent to all peers
 - Active leader’s registered state recovery function provides “blob”
 - Data returned to recovering app
- Recovery time
 - Few msec to restart process
 - State recovery limited by time to transfer blob between procs

Example: BGP

Simulator layout

- Multiple BGP siblings
 - All siblings store state locally in a routing table
 - Track sequence number of last update processed into table
 - Not really executing BGP algo, just processing faux updates into table
- Multiple shims
 - Each receives messages from faux external BGPs
 - Shim leader controls return channel
 - Relay messages to siblings
 - Retain local ring buffer of last N messages

Example: BGP

- Sibling restarts
 - Checks incarnation>0
 - Requests state recovery data
- Active sibling leader
 - Receives callback with request
 - Binary copies state table into return message, along with sequence number of last update
 - Orcm returns resulting blob to caller
- Restarted sibling
 - Places state table blob into its table
 - Notifies shims “ready” as of given sequence number
 - Shims send catch-up updates, continue from that point with normal updates

Node Replacement/Addition

- Failure or shutdown
 - Failure announced to all processes to permit leadership switch, if required
 - Orcm-sched relocates procs up to some max #times (can be unlimited)
- Auto-boot of local daemon on power up
 - Daemon announces to DVM
 - Orcm-sched adds node to available resources
- Processes will map to new resource as start/restart demands
 - Future: rebalance existing load upon node availability

Current Status

- Overall system is functional, tested on router and cluster
 - Expected scaling issues (UDP stack limits on total messaging rate) observed
 - Several potential paths to improve scalability have been identified
 - Sender-based TCP messaging successfully avoids scaling problems
- No further development underway

Concluding Remarks



<http://www.open-mpi.org>
<http://www.open-mpi.org/projects/orcm>